

Terraform Essential Components

Whether you're just starting with Terraform or looking to brush up on the basics, this cheatsheet optimizes your workflow and helps you prepare for the certification.

Provider

- Plugins that allow Terraform to interact with specific infrastructure resources.
- Act as an interface between Terraform and the underlying infrastructure translating the Terraform configuration into the appropriate API calls
- Each provider has its own set of resources and data sources
- Providers exist even for applications that are not Cloud Vendor specific

```
provider "aws" {
  region = "us-east-1"
}

provider "aws" {
  alias   = "euwest"
  region = "eu-west-1"
}

provider "azurem" {
  features {}
}

provider "kubernetes" {
  config_path = "~/.kube/config"
}
```

- You can define multiple providers of the same type by using the alias key word, in resources or datasources, you will have to specify either:

- Resource: a provider parameter (example below)
- Module: a provider block (example below)

```
# Alias Examples for Resources and Modules
resource "aws_instance" "this" {
  provider = aws.euwest
  ami      = data.aws_ami.ubuntu.id
  instance_type = "t3.micro"
}

module "instance" {
  source = "../instance"
  providers = {
    aws = aws.euwest
  }
}
```

Resource

- Components of infrastructure that Terraform is able to manage: vms, vpcs, pods
- Each resource has a type, a name, some configurable arguments and some exposed attributes
- Terraform handles dependencies by default, but in order to create these dependencies between resources, you need to link the resources between them by using **type.name.attribute**

```
resource "azurerm_resource_group" "this" {
  name     = "rg1"
  location = "West Europe"
}

resource "azurerm_network_security_group" "this" {
  name     = "sg1"
  name     = azurerm_resource_group.this.location # ensures rg created before nsg and links to the above rg
  resource_group_name = azurerm_resource_group.this.name # ensures rg created before nsg and links to the above rg
}
```

- How to handle the documentation
 - **Example Usage** → How to use the resource
 - **Argument Reference** → What you can configure
 - **Attribute Reference** → What is the resource exposing
 - **Import** → What you can import

Data Source

- Used to get different information that was created outside of Terraform
- Every provider has its own set of data sources
- Can be used inside of resources/locals/outputs
- Referenced with **data.type.name.attribute**

```
data "aws_ami" "ubuntu" {
  most_recent = true

  filter {
    name     = "name"
    values   = ["ubuntu*"]
  }
}

resource "aws_instance" "this" {
  ami          = data.aws_ami.ubuntu.id
  instance_type = "t3.micro"
}
```

Output

- Used for exposing different attributes of a resource, variable, data source or local
- Exist outside of providers
- In Terraform Modules, they are used to expose some attributes that can be used in other configurations (either modules, resources, datasources, locals, providers)
- Supports three arguments, value (required), description and sensitive (whether or not the output is sensitive)

```
locals {
  a_list = [for i in range(10) : i]
}

output "a_list" {
  description = "This will print the elements of a_list"
  value       = local.a_list
}

# a_list = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Variable

- Declared with variable, referenced with var
- Simple types: number, string, bool
- Complex types: list, map, object, set → built from simple types
- Null types: used for omission
- 3 essential parameters when configuring a variable, all optional though: type, description and default (value)
- 3 other parameters: validation, sensitive, nullable

```
variable "a_string" {
  description = "This is a string"
  type        = string
  default     = "a_string"
}

variable "a_number" {
  description = "This is a number"
  type        = number
  default     = 10
}

variable "a_bool" {
  description = "This is a bool"
  type        = bool
  default     = false
}

variable "a_list" {
  description = "This is a list"
  type        = list(string)
  default     = ["elem1", "elem2"]
}
```

- Taking into account the definition precedence, you can assign values to variables:
 - Inside of them
 - In an environment variable
 - In a terraform.tfvars file
 - In a *.auto.tfvars file where * can be any name you want
 - Using the -var option or -var-file option

- Terraform will use the last value it finds when you are declaring a variable value in multiple sources, but you will not be able to declare the same variable in the same source more than once, as this will result in an error

```
variable "a_map" {
  description = "This is a map"
  type        = map(string)
  default = {
    elem1 = "value1"
  }
}

variable "an_object" {
  description = "This is an object"
  type        = object({
    name = string
    power = number
  })
  default = {
    name = "elem1"
    power = 5
  }
}
```

Local Variable

- Similar to a variable, but used to assign a name to an expression
- Exists in a locals block
- You can use constant values, attributes from resource, dynamically generate expression with loops or conditional statements

```
locals {
  expression1 = format("Hello %s!", "Mike")
  ternary_expression = 3 > 2 ? 1 : 4 # If 3 is greater than 2 which it is, it will use value 1, otherwise value 4
  for_expression_list = [for i in range(5) : i] # This will create a list of numbers from 0 to 4
  for_expression_map = { for i in range(10) : format("Number_%d", i) => i } # This will create a map with key value pairs in the format Number_0 => 0
}
```

Provisioner

- Used for doing operations outside Terraform's lifecycle
- Exist only inside of a resource, there is also a special type of resource called null_resource that can be leveraged to work with providers (these resources aren't doing anything on their own)
- Support **when** (this refers to when do you want them to be applied **on_create** which is default or **on_destroy**) and **on_failure** (which refers to the behavior they have when they encounter an error: **continue** to ignore the error, **fail** which is the default to behave exactly like any other statement in Terraform)
- 3 available provisioners, local-exec (runs a local command), remote-exec (runs a remote command), file (copies a file from local to remote)
- Remote provisioners need a connection block to identify where they should connect
- Should be used only as a last resort - you should use configuration management tools (Ansible, Chef, Puppet) to achieve these types of tasks

```
resource "null_resource" "this" {
  provisioner "local-exec" {
    command = "ls -l"
  }

  connection {
    type        = "ssh"
    user        = "ubuntu"
    private_key = file("~/.ssh/id_rsa")
    host        = "your_host"
  }

  provisioner "remote-exec" {
    inline = [
      "echo hello"
    ]
  }

  provisioner "file" {
    source      = "../file.yaml"
    destination = "../file.yaml"
  }
}
```